

t-SNE-Balanced-AFF-Review

November 22, 2018

1 t-SNE Visualization on Amazon Food Review Dataset

1.1 Import Required Modules

```
In [1]: import os # for file management
import shutil # for file management
from pathlib import Path # for file management
import sqlite3
import pandas as pd
import numpy as np
import csv
from tqdm import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
import time # for time measurement
import imageio # for GIF creation

from sklearn.feature_extraction.text import CountVectorizer # for Bag Of Words
from sklearn.feature_extraction.text import TfidfVectorizer # for text to vector creation
from gensim.models import Word2Vec

from sklearn.preprocessing import StandardScaler # for Column Standardization - DO WE NEED THIS?
from sklearn.manifold import TSNE # for t-SNE

In [2]: ## Configure Matplotlib for nice image in PDF
from IPython.display import set_matplotlib_formats
set_matplotlib_formats('pdf', 'png')
plt.rcParams['savefig.dpi'] = 75
plt.rcParams['figure.figsize'] = 10,6
plt.rcParams['axes.labelsize'] = 18
plt.rcParams['axes.titlesize'] = 20
plt.rcParams['font.size'] = 10
plt.rcParams['lines.linewidth'] = 2.0
plt.rcParams['lines.markersize'] = 8

In [3]: # All the outputs generated by this notebook will be placed in separate folder
output_dir = 'Output'
if not os.path.exists(output_dir):
    os.makedirs(output_dir)
```

1.2 Load Data

```
In [4]: # Loading all the reviews from the database
# cleaned.sqlite already has reviews of 1,2,4,5 ratings only (also changed to +ive/-ive)
con = sqlite3.connect('./cleaned.sqlite')

df = pd.read_sql_query("""SELECT * from Reviews""", con)
df.head()
```

```
Out[4]:
```

	index	Id	ProductId	UserId	ProfileName	\
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	
1	138688	150506	0006641040	A2IW4PEEK02ROU	Tracy	
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	
3	138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"	
4	138691	150509	0006641040	A3CMRKGE0P909G	Teresa	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	0	0	1	939340800	
1	1	1	1	1194739200	
2	1	1	1	1191456000	
3	1	1	1	1076025600	
4	3	4	1	1018396800	

	Summary	\
0	EVERY book is educational	
1	Love the book, miss the hard cover version	

```

2         chicken soup with rice months
3     a good swingy rhythm for reading aloud
4         A great way to learn the months

                                Text \
0  this witty little book makes my son laugh at l...
1  I grew up reading these Sendak books, and watc...
2  This is a fun way for children to learn their ...
3  This is a great little book to read aloud- it ...
4  This is a book of poetry about the months of t...

                                CleanedText
0  b'witti littl book make son laugh loud recit c...
1  b'grew read sendak book watch realli rosi movi...
2  b'fun way children learn month year learn poem...
3  b'great littl book read nice rhythm well good ...
4  b'book poetri month year goe month cute littl ...

```

```
In [5]: df.describe()
```

```

Out[5]:
count  364106.000000  364106.000000  364106.000000  \
mean    261221.056821  282777.564772    1.738411
std     152361.122483  164601.735167    6.716471
min         0.000000    1.000000    0.000000
25%     129625.250000  140699.250000    0.000000
50%     257307.500000  278947.500000    0.000000
75%     396338.750000  428557.750000    2.000000
max     525813.000000  568454.000000   866.000000

                                HelpfulnessDenominator  Score  Time
count          364106.000000  364106.000000  3.641060e+05
mean              2.186231    0.843164  1.296157e+09
std              7.339767    0.363647  4.859821e+07
min              0.000000    0.000000  9.393408e+08
25%              0.000000    1.000000  1.270858e+09
50%              1.000000    1.000000  1.311379e+09
75%              2.000000    1.000000  1.332893e+09
max             878.000000    1.000000  1.351210e+09

```

```
In [6]: df.dtypes
```

```

Out[6]:
index          int64
Id             int64
ProductId      object
UserId        object
ProfileName    object
HelpfulnessNumerator  int64
HelpfulnessDenominator  int64
Score          int64
Time           int64
Summary        object
Text           object
CleanedText    object
dtype: object

```

```

In [7]: # Split data
# positive review score, negative review score and review text as seperate dataframes
df_text = df['CleanedText']
print(df_text.shape)
df_text.head()

```

```
(364106,)
```

```

Out[7]:
0  b'witti littl book make son laugh loud recit c...
1  b'grew read sendak book watch realli rosi movi...
2  b'fun way children learn month year learn poem...
3  b'great littl book read nice rhythm well good ...
4  b'book poetri month year goe month cute littl ...
Name: CleanedText, dtype: object

```

```

In [8]: # Procedure of t-SNE generation is same for all types of vector analysis
# So creating a common function to generate t-SNE visualization
# using passed values

def genTSNEGif(std_data, ndp, p, itr_list, file_prefix, closePlt=False):

```

```

'''
Fuction which genrate t-SNE visualtion for each itr_list using given ndp and p
Generates a GIF and stores it under '{img_name}.gif'
Where:
    std_data - Column Standardized Data
    ndp - Number of Data Points to consider in std_data
    p - Perplexity
    itr_list - List of iterations, each iteration will be a frame in GIF
    file_prefix - Prefix to the name of GIF image
    closePlt - If you do not want to display the generated image in Notebook
'''
image_name = '{0}_tsne_ndp_{1}_p_{2}.gif'.format(file_prefix,ndp,p)
print('No.Of Data Points - {0}, Perplexity - {1}, Iterations - {2}, ImageName - {3}'.format(
    ndp, p, itr_list, image_name))

# list to hold the frames
frames = []
p_data = std_data
p_labels = final_reviews_scores[0:ndp]

#print('t-SNE Data Points {0} and its Labels {1}'.format(p_data.shape, p_labels.shape))
for itr_val in itr_list:
    img_title = '{0}-ndp={1} p={2} itr={3}'.format(file_prefix, ndp, p, itr_val)

    time_start = time.time()

    model = TSNE(n_components=2,random_state=0,perplexity=p,n_iter=itr_val) # ,verbose=2
    tsne_data = model.fit_transform(p_data)
    time_elapsed = time.time() - time_start
    print('{0} ==> t-SNE done! Time elapsed: {1} seconds'.format(img_title, time.time() - time_start))

    tsne_data = np.vstack((tsne_data.T,p_labels)).T
    #print(tsne_data.shape)
    #tsne_data[:4]
    tsne_df = pd.DataFrame(tsne_data,columns=['Dim_1','Dim_2','Score'])
    #tsne_df.head()
    g = sns.FacetGrid(tsne_df,hue='Score',height=10).map(plt.scatter, 'Dim_1', 'Dim_2').add_legend();
    g.fig.suptitle(img_title);
    g.fig.canvas.draw();
    image = np.frombuffer(g.fig.canvas.tostring_rgb(), dtype='uint8')
    image = image.reshape(g.fig.canvas.get_width_height()[::-1] + (3,))
    frames.append(image)

    if (closePlt == True):
        plt.close()

kwargs_write = {'fps':1.0, 'quantizer':'nq'}
imageio.mimsave(Path.cwd() / output_dir / image_name, frames, fps=1)

return

```

1.3 Training Data for Visualization - 2K Points

```

In [9]: # we can't process all 364K reviews, selecting a subset of it
total_data_set_size = 2000

# Create a Balanced dataset having both +ive and -ive reviews
df_positive_reviews = df[df.Score == 1].sample(int(total_data_set_size/2))
df_negative_reviews = df[df.Score == 0].sample(int(total_data_set_size/2))

final_reviews = pd.concat([df_positive_reviews, df_negative_reviews])
final_reviews_scores = final_reviews['Score']

print('Shape of Training Data {0}'.format(final_reviews.shape))
print('Shape of Training Label {0}'.format(final_reviews_scores.shape))

```

Shape of Training Data (2000, 12)
Shape of Training Label (2000,)

In [10]: final_reviews.head()

```

Out[10]:
   index   Id  ProductId  UserId \
160283  202154  219052  B001CRAWCQ  A15C0QFMZ6JTE1
350384   44522   48450  B006H34CUS  A1C935N5QX5S4S
32515   178397  193447  B000BBDZ82  A3CZNNTEY657TE
60680    17311   18879  B000F4H5FY  A2434W0L2ZL5X1

```

43145	381893	412923	B000E4C2LW	A3TXD09392M8NJ	
-------	--------	--------	------------	----------------	--

	ProfileName	HelpfulnessNumerator	\
160283	Ivana Pandurovic	1	
350384	Jamie Radcliff	0	
32515	C. Gardner "SgtCheeseNOLS"	1	
60680	Brian	0	
43145	Beth Mitchum "Ultravioletlove.com"	1	

	HelpfulnessDenominator	Score	Time	\
160283	2	1	1303344000	
350384	0	1	1346716800	
32515	1	1	1325462400	
60680	0	1	1291507200	
43145	1	1	1275350400	

	Summary	\
160283	LOVE Illy cafecito	
350384	Not too bad	
32515	LOVE IT!	
60680	Pretty good! Great Price!	
43145	Good Stuff at a Great Price	

	Text	\
160283	i am european, addicted not only to coffee but...	
350384	I received a sample of this from Influenster a...	
32515	I first tried this tea at my day spa I go to f...	
60680	I am aware that tea bags do not give the same ...	
43145	I know what you're thinking. Corn flakes are ...	

	CleanedText
160283	b'european addict coffe good coffe illi best c...
350384	b'receiv sampl influenst pretti excit tri inte...
32515	b'first tri tea day spa massag delici tea comp...
60680	b'awar tea bag give flavour looseleaf stuff st...
43145	b'know your think corn flake corn flake right ...

2 Bag of Words (BoW)

```
In [11]: # Create Vectors
count_vect = CountVectorizer(ngram_range=(1,2)) # create an instance
final_counts = count_vect.fit_transform(final_reviews['CleanedText'].values)
print('Shape of BoW Vectorizer: ', final_counts.get_shape())
print('Total no.of unique words: ', final_counts.get_shape()[1])


# Standardize the Data
standardized_data = StandardScaler().fit_transform(final_counts.toarray().astype(np.float64)) #, with_mean=False
print('Shape of Standardized data', standardized_data.shape)

Shape of BoW Vectorizer: (2000, 66148)
Total no.of unique words: 66148
Shape of Standardized data (2000, 66148)

In [12]: # Generating t-SNE for BoW
genTSNEgif(standardized_data, len(standardized_data), 30, range(3000,3001,1000), 'BoW-std')

# For various iterations I use below looping
#dense_mat = final_counts.toarray().astype(np.float64)
#for p in range(10, 101, 10):
#    genTSNEgif(dense_mat, len(dense_mat), p, range(1000,5001,1000), 'BoW',closePlt=True)

No.Of Data Points - 2000, Perplexity - 30, Iterations - range(3000, 3001, 1000), ImageName - BoW-std_tsne_ndp_2000_p_30.gif
BoW-std-ndp=2000 p=30 itr=3000 ==> t-SNE done! Time elapsed: 598.4183218479156 seconds
```



3 TFIDF

In [13]: # Create Vectors

```
tf_idf_vec = TfidfVectorizer(ngram_range=(1,2))
final_counts = tf_idf_vec.fit_transform(final_reviews['CleanedText']).values

#.fit_transform(final_reviews['CleanedText'].values)
print('Shape of BoW Vectorizer: ', final_counts.get_shape())
print('Total no.of unique words: ', final_counts.get_shape()[1])


# Standardize the Data
standardized_data = StandardScaler().fit_transform(final_counts.toarray().astype(np.float64)) #, with_mean=False
print('Shape of Standardized data', standardized_data.shape)
```

```
Shape of BoW Vectorizer: (2000, 66148)
Total no.of unique words: 66148
Shape of Standardized data (2000, 66148)
```

```
In [14]: genTSNEGif(standardized_data, len(standardized_data), 30, range(3000,3001,1000), 'tfidf-std')
```

```
# For finding optimal perplexity and iteration, I execute below loop
#dense_mat = final_counts.toarray().astype(np.float64)
#for p in range(10, 61, 10):
#    genTSNEGif(dense_mat, len(dense_mat), p, range(1000,6001,1000), 'tfidf',closePlt=True)
```

No.Of Data Points - 2000, Perplexity - 30, Iterations - range(3000, 3001, 1000), ImageName - tfidf-std_tsne_ndp_2000_p_30.gif
 tfidf-std-ndp=2000 p=30 itr=3000 ==> t-SNE done! Time elapsed: 599.3038096427917 seconds



4 Word2Vec

I am creating vectors having 50 dimensions. Just a random value, not inherent calculation I made on this size decision.

```
In [15]: # Create List array for creating own W2V
list_of_sent = []
for sent in final_reviews['CleanedText'].values:
    list_of_sent.append(sent.decode("utf-8").split())
```

```
print(final_reviews.CleanedText.values[0])
print(len(list_of_sent), list_of_sent[0])
```

```
b'europaean addict coffe good coffe illi best cant start day without illi espresso doubl one'
2000 ['europaean', 'addict', 'coffe', 'good', 'coffe', 'illi', 'best', 'cant', 'start', 'day', 'without', 'illi', 'espresso', 'doubl', 'one']
```

```
In [16]: # Required dimension
w2v_d = 50
```

```
# Considering words that are occurred atleast 5 times in the corpus
w2v_model = Word2Vec(list_of_sent, min_count=5, size=w2v_d, workers=4)

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times : ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times : 2017
sample words ['addict', 'coffe', 'good', 'illi', 'best', 'cant', 'start', 'day', 'without', 'espresso', 'doubl', 'one', 'receiv', 'sampl', '']
```

4.1 Avg-W2V

```
In [17]: # Computing average w2v for each review in selected training dataset
review_vectors = []
for sent in tqdm(list_of_sent, ascii=True):
    sent_vec = np.zeros(w2v_d) # array to hold the vectors. Initially assuming no vectors in this review
    no_of_words_in_review = 0 # number of words with valid vector in this review

    # count all the words (that are in w2v model) and take average
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            no_of_words_in_review += 1
    if no_of_words_in_review != 0:
        sent_vec /= no_of_words_in_review
    review_vectors.append(sent_vec)

print(len(review_vectors))
print(len(review_vectors[0]))
```

```
100%|#####| 2000/2000 [00:02<00:00, 733.21it/s]
```


```
2000
50
```

```
In [18]: # t-SNE using Average Word2Vec
```

```
genTSNEGif(review_vectors, len(review_vectors), 70, range(3000,3001,1000), 'avg-w2v')

# for finding perplexity and iteration, I use below loop
# for p in range(10, 101, 10):
#     genTSNEGif(review_vectors, len(review_vectors), p, range(1000,5001,1000), 'avg-w2v', closePlt=True)
```

```
No.Of Data Points - 2000, Perplexity - 70, Iterations - range(3000, 3001, 1000), ImageName - avg-w2v-tsne-ndp-2000_p-70.gif
avg-w2v-ndp=2000 p=70 itr=3000 ==> t-SNE done! Time elapsed: 56.379451274871826 seconds
```



4.2 TFIDF Weighted W2V

Computing tfidf weighted w2v over the selected training dataset

```
In [19]: # Create tf-idf vector matrix
tf_idf_model = TfidfVectorizer(ngram_range=(1,2))
tf_idf_matrix = tf_idf_model.fit_transform(final_reviews['CleanedText'].values)

# Create dictionary having words (features) as keys, its tf-idf values as values
tf_idf_dict = dict(zip(tf_idf_model.get_feature_names(), list(tf_idf_model.idf_)))
len(tf_idf_dict)
```

Out[19]: 66148

```
In [20]: tf_idf_feat = tf_idf_model.get_feature_names()

# Computing tf-idf weighted w2v for each review in selected training dataset
review_vectors = []
for sent in tqdm(list_of_sent, ascii=True):
    sent_vec = np.zeros(w2v_d) # array to hold the vectors
    no_of_words_in_review = 0 # number of words with valid vector in this review
```



```

# count all the words (that are in w2v model) and take average
for word in sent:
    if word in w2v_words:
        vec = w2v_model.wv[word]
        # calculate tf-idf weighted w2v value for this word
        tf_idf = tf_idf_dict[word] * (sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        no_of_words_in_review += 1
if no_of_words_in_review != 0:
    sent_vec /= no_of_words_in_review
review_vectors.append(sent_vec)

print(len(review_vectors))
print(len(review_vectors[0]))

```

100%|#####| 2000/2000 [00:03<00:00, 563.31it/s]

2000
50

In [21]: # t-SNE using tf-idf weighted s2v


```

genTSNEgif(review_vectors, len(review_vectors), 80, range(3000,3001,1000), 'tfidf-weighted-w2v')

#for p in range(10, 101, 10):
#    genTSNEgif(review_vectors, len(review_vectors), p, range(1000,5001,1000), 'tfidf-weighted-w2v',closePlt=True)

```

No.Of Data Points - 2000, Perplexity - 80, Iterations - range(3000, 3001, 1000), ImageName - tfidf-weighted-w2v-tsne_ndp_2000_p_80.gif
tfidf-weighted-w2v-ndp=2000 p=80 itr=3000 ==> t-SNE done! Time elapsed: 52.13369131088257 seconds



5 t-SNE Observation

Due to memory limitation, only 2000 reviews out of 364K reviews of Amazon Fine Food Review has been taken

Created a Balanced dataset having 1000 positive and negative reviews respectively.

Generated t-SNE 2-D visualization from the vectors created with 1) Bag of Words, 2) TF-IDF, 3) Average Word 2 Vec and 4) TF-IDF Weighter Word2Vec.

I have generated my own Word2Vec from those 2K reviews having dimensionve of 50

- Both BoW and TF-IDF generated a small circled image, having both Positive and Negative reviews overlapped in same place
 - There are some outliers kind of observation as well seen
- In Avg-W2V, U-Shaped curve is seen. But still here as well both postivie and negative reviews overlap each other
- In TFIDF Weighter W2V, kind of sign wave observed, but still both positive and negative reviews overlap occurred.

In conclusion, no clustering seen from t-SNE. So we can't have a linearly seperable classification function to classify that whether a review is positive or not.